

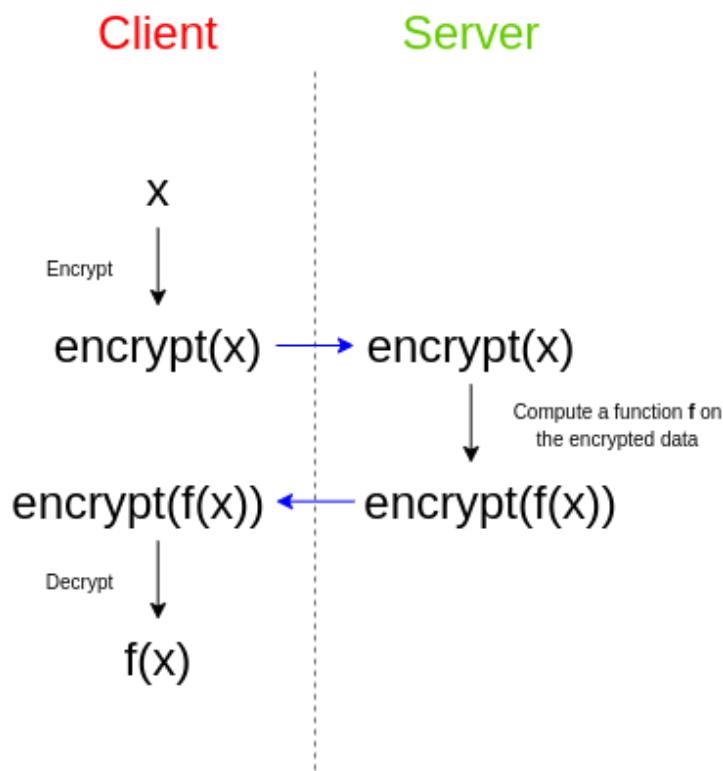
Exploring Fully Homomorphic Encryption

2020 Jul 20

[See all posts](#)

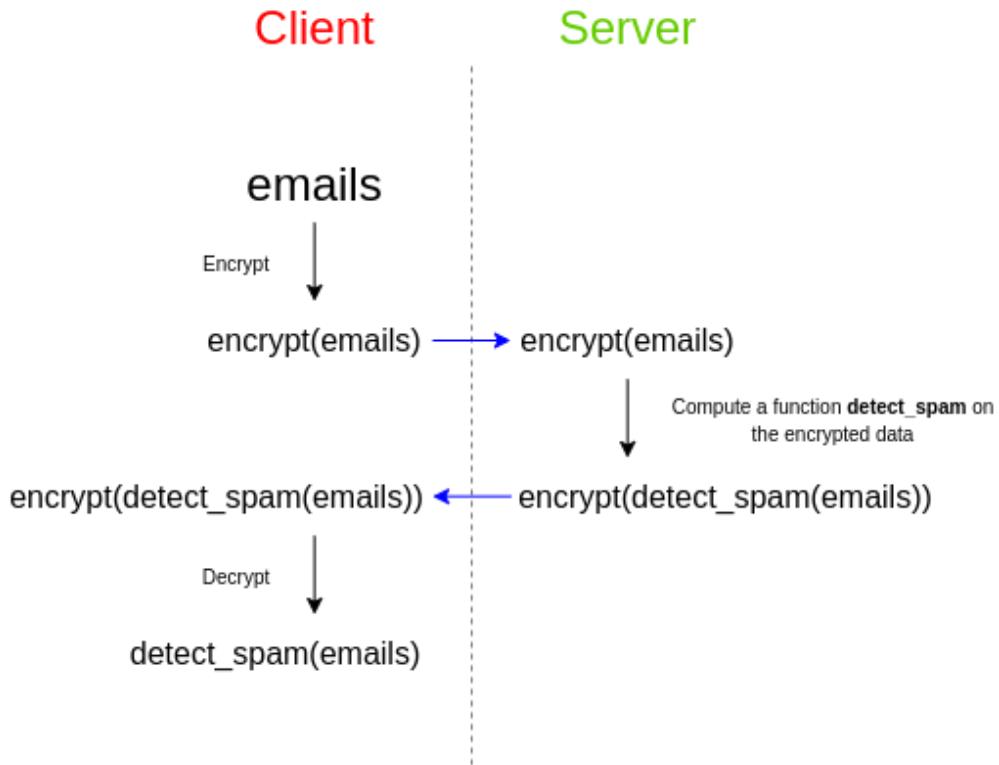
Special thanks to Karl Floersch and Dankrad Feist for review

Fully homomorphic encryption has for a long time been considered one of the holy grails of cryptography. The promise of fully homomorphic encryption (FHE) is powerful: it is a type of encryption that allows a third party to perform computations on encrypted data, and get an encrypted result that they can hand back to whoever has the decryption key for the original data, *without* the third party being able to decrypt the data or the result themselves.



As a simple example, imagine that you have a set of emails, and you want to use a third party spam filter to check whether or not they are spam. The spam filter has a desire for *privacy of*

their algorithm: either the spam filter provider wants to keep their source code closed, or the spam filter depends on a very large database that they do not want to reveal publicly as that would make attacking easier, or both. However, you care about the *privacy of your data*, and don't want to upload your unencrypted emails to a third party. So here's how you do it:



Fully homomorphic encryption has many applications, including in the blockchain space. One key example is that can be used to implement privacy-preserving light clients (the light client hands the server an encrypted index i , the server computes and returns $\text{data}[0] * (i = 0) + \text{data}[1] * (i = 1) + \dots + \text{data}[n] * (i = n)$, where $\text{data}[i]$ is the i 'th piece of data in a block or state along with its Merkle branch and $(i = k)$ is an expression that returns 1 if $i = k$ and otherwise 0; the light client gets the data it needs and the server learns nothing about what the light client asked).

It can also be used for:

- More efficient **stealth address protocols**, and more generally scalability solutions to privacy-preserving protocols that today require each user to personally scan the entire blockchain for incoming transactions
- Privacy-preserving data-sharing marketplaces that let users allow some specific computation to be performed on their data while keeping full control of their data for themselves

- An ingredient in more powerful cryptographic primitives, such as more efficient multi-party computation protocols and perhaps eventually obfuscation

And it turns out that fully homomorphic encryption is, conceptually, not that difficult to understand!

Partially, Somewhat, Fully homomorphic encryption

First, a note on definitions. There are different kinds of homomorphic encryption, some more powerful than others, and they are separated by what kinds of functions one can compute on the encrypted data.

- **Partially homomorphic encryption** allows evaluating only a *very limited* set of operations on encrypted data: either just additions (so given `encrypt(a)` and `encrypt(b)` you can compute `encrypt(a+b)`), or just multiplications (given `encrypt(a)` and `encrypt(b)` you can compute `encrypt(a*b)`).
- **Somewhat homomorphic encryption** allows computing additions as well as a *limited* number of multiplications (alternatively, polynomials up to a limited degree). That is, if you get `encrypt(x1) ... encrypt(xn)` (assuming these are "original" encryptions and not already the result of homomorphic computation), you can compute `encrypt(p(x1 ... xn))`, as long as $p(x_1 \dots x_n)$ is a polynomial with degree $< D$ for some specific degree bound D (D is usually very low, think 5-15).
- **Fully homomorphic encryption** allows unlimited additions and multiplications. Additions and multiplications let you replicate any binary circuit gates ($\text{AND}(x, y) = x*y$, $\text{OR}(x, y) = x+y-x*y$, $\text{XOR}(x, y) = x+y-2*x*y$ or just $x+y$ if you only care about even vs odd, $\text{NOT}(x) = 1-x$...), so this is sufficient to do arbitrary computation on encrypted data.

Partially homomorphic encryption is fairly easy; eg. RSA has a multiplicative homomorphism: $\text{enc}(x) = x^e$, $\text{enc}(y) = y^e$, so $\text{enc}(x) * \text{enc}(y) = (xy)^e = \text{enc}(xy)$. Elliptic curves can offer similar properties with addition. Allowing *both* addition and multiplication is, it turns out, significantly harder.

A simple somewhat-HE algorithm

Here, we will go through a somewhat-homomorphic encryption algorithm (ie. one that supports a limited number of multiplications) that is surprisingly simple. A more complex version of this category of technique was used by Craig Gentry to create [the first-ever fully](#)

homomorphic scheme in 2009. More recent efforts have switched to using different schemes based on vectors and matrices, but we will still go through this technique first.

We will describe all of these encryption schemes as *secret-key* schemes; that is, the same key is used to encrypt and decrypt. Any secret-key HE scheme can be turned into a public key scheme easily: a "public key" is typically just a set of many encryptions of zero, as well as an encryption of one (and possibly more powers of two). To encrypt a value, generate it by adding together the appropriate subset of the non-zero encryptions, and then adding a random subset of the encryptions of zero to "randomize" the ciphertext and make it infeasible to tell what it represents.

The secret key here is a large prime, p (think of p as having hundreds or even thousands of digits). The scheme can only encrypt 0 or 1, and "addition" becomes XOR, ie. $1 + 1 = 0$. To encrypt a value m (which is either 0 or 1), generate a large random value R (this will typically be even larger than p) and a smaller random value r (typically much smaller than p), and output:

$$\text{enc}(m) = R * p + r * 2 + m$$

To decrypt a ciphertext ct , compute:

$$\text{dec}(ct) = (ct \bmod p) \bmod 2$$

To add two ciphertexts ct_1 and ct_2 , you simply, well, add them: $ct_1 + ct_2$. And to multiply two ciphertexts, you once again... multiply them: $ct_1 * ct_2$. We can prove the homomorphic property (that the sum of the encryptions is an encryption of the sum, and likewise for products) as follows.

Let:

$$ct_1 = R_1 * p + r_1 * 2 + m_1$$

$$ct_2 = R_2 * p + r_2 * 2 + m_2$$

We add:

$$ct_1 + ct_2 = R_1 * p + R_2 * p + r_1 * 2 + r_2 * 2 + m_1 + m_2$$

Which can be rewritten as:

$$(R_1 + R_2) * p + (r_1 + r_2) * 2 + (m_1 + m_2)$$

Which is of the exact same "form" as a ciphertext of $m_1 + m_2$. If you decrypt it, the first $\text{mod } p$ removes the first term, the second $\text{mod } 2$ removes the second term, and what's left is $m_1 + m_2$ (remember that if $m_1 = 1$ and $m_2 = 1$ then the 2 will get absorbed into the second term and you'll be left with zero). And so, voila, we have additive homomorphism!

Now let's check multiplication:

$$ct_1 * ct_2 = (R_1 * p + r_1 * 2 + m_1) * (R_2 * p + r_2 * 2 + m_2)$$

Or:

$$\begin{aligned} & (R_1 * R_2 * p + r_1 * 2 + m_1 + r_2 * 2 + m_2) * p + \\ & (r_1 * r_2 * 2 + r_1 * m_2 + r_2 * m_1) * 2 + \\ & (m_1 * m_2) \end{aligned}$$

This was simply a matter of expanding the product above, and grouping together all the terms that contain p , then all the remaining terms that contain 2, and finally the remaining term which is the product of the messages. If you decrypt, then once again the $\text{mod } p$ removes the first group, the $\text{mod } 2$ removes the second group, and only $m_1 * m_2$ is left.

But there are two problems here: first, the size of the ciphertext itself grows (the length roughly doubles when you multiply), and second, the "noise" (also often called "error") in the smaller $\text{mod } 2$ term also gets quadratically bigger. Adding this error into the ciphertexts was necessary because the security of this scheme is based on the approximate GCD problem:

Approximate GCD problem

In the **approximate GCD problem** you are given n "near" multiples $\{m_1, m_2, \dots, m_n\}$ of some unknown positive integer p , i.e.

$$m_1 = p \cdot q_1 + r_1, \quad 0 \leq r_1 < p,$$

$$m_2 = p \cdot q_2 + r_2, \quad 0 \leq r_2 < p,$$

...

$$m_n = p \cdot q_n + r_n, \quad 0 \leq r_n < p,$$

and you need to find out p , without the knowledge of any q_i and r_i .

Had we instead used the "exact GCD problem", breaking the system would be easy: if you just had a set of expressions of the form $p * R_1 + m_1, p * R_2 + m_2, \dots$, then you could use the Euclidean algorithm to efficiently compute the greatest common divisor p . But if the

ciphertexts are only *approximate* multiples of p with some "error", then extracting p quickly becomes impractical, and so the scheme can be secure.

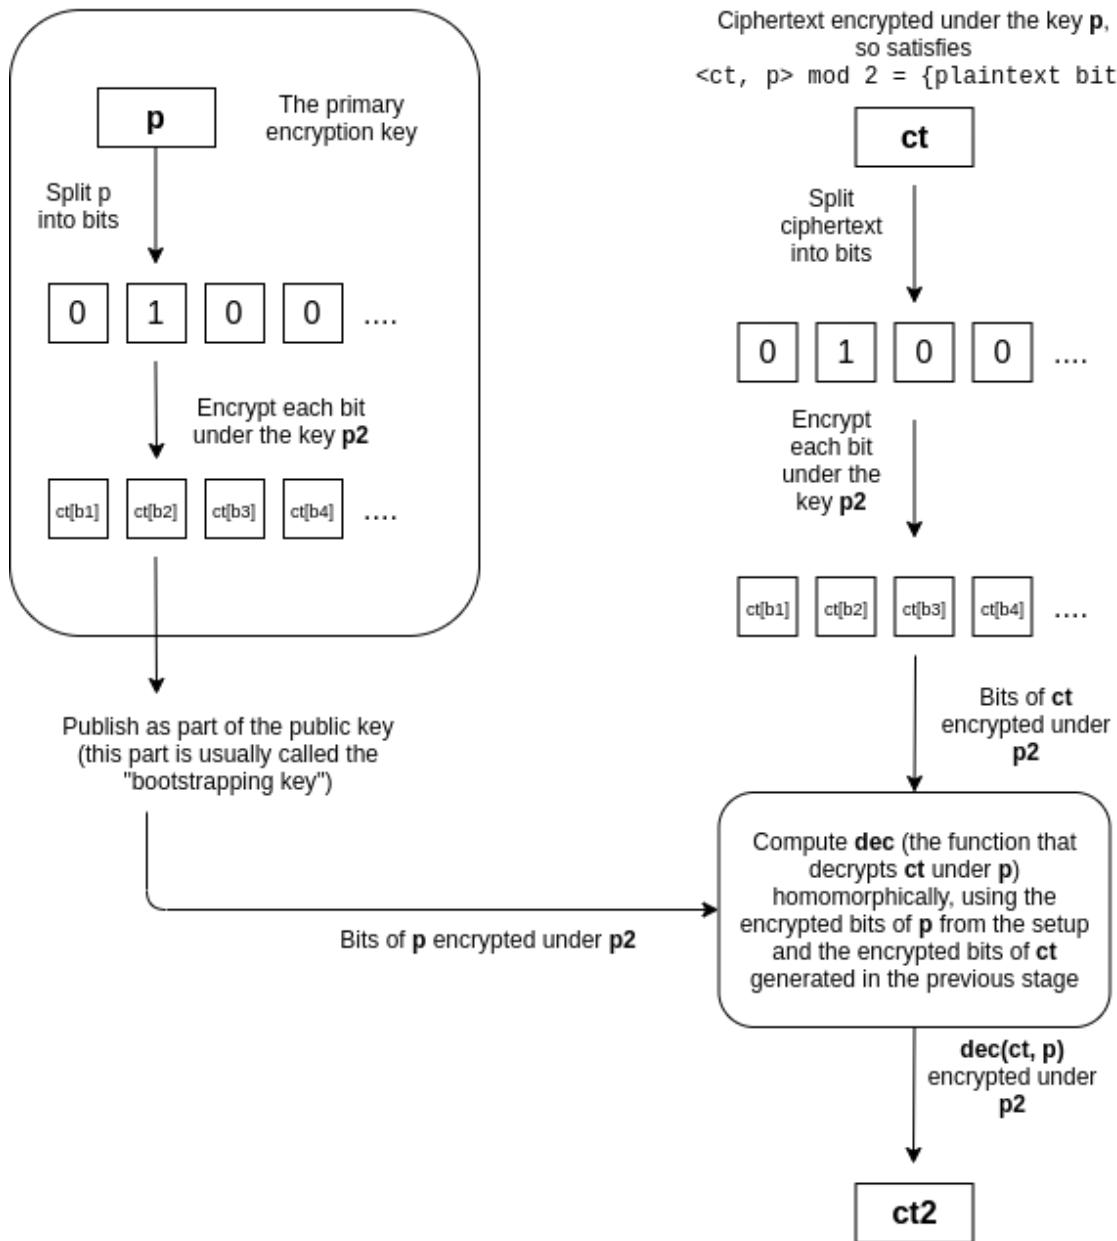
Unfortunately, the error introduces the inherent limitation that if you multiply the ciphertexts by each other enough times, the error eventually grows big enough that it exceeds p , and at that point the **mod p** and **mod 2** steps "interfere" with each other, making the data unextractable. This will be an inherent tradeoff in all of these homomorphic encryption schemes: extracting information from *approximate* equations "with errors" is much harder than extracting information from exact equations, but any error you add quickly increases as you do computations on encrypted data, bounding the amount of computation that you can do before the error becomes overwhelming. And **this is why these schemes are only "somewhat" homomorphic.**

Bootstrapping

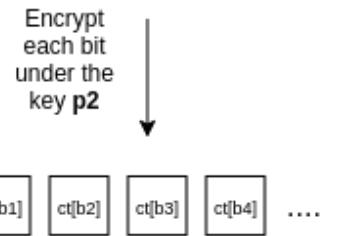
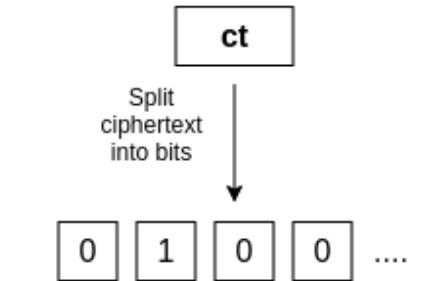
There are two classes of solution to this problem. First, in many somewhat homomorphic encryption schemes, there are clever tricks to make multiplication only increase the error by a constant factor (eg. 1000x) instead of squaring it. Increasing the error by 1000x still sounds by a lot, but keep in mind that if p (or its equivalent in other schemes) is a 300-digit number, that means that you can multiply numbers by each other 100 times, which is enough to compute a very wide class of computations. Second, there is Craig Gentry's technique of "bootstrapping".

Suppose that you have a ciphertext ct that is an encryption of some m under a key p , that has a lot of error. The idea is that we "refresh" the ciphertext by turning it into a new ciphertext of m under another key p_2 , where this process "clears out" the old error (though it will introduce a fixed amount of new error). The trick is quite clever. The holder of p and p_2 provides a "bootstrapping key" that consists of an encryption of *the bits of p* under the key p_2 , as well as the public key for p_2 . Whoever is doing computations on data encrypted under p would then take the bits of the ciphertext ct , and individually encrypt these bits under p_2 . They would then *homomorphically compute the decryption under p* using these ciphertexts, and get out the single bit, which would be m encrypted under p_2 .

Setup



Ciphertext encrypted under the key p ,
so satisfies
 $\langle ct, p \rangle \bmod 2 = \{\text{plaintext bit}\}$



Compute dec (the function that
decrypts ct under p)
homomorphically, using the
encrypted bits of p from the setup
and the encrypted bits of ct
generated in the previous stage

$\text{dec}(ct, p)$
encrypted under
 p_2

$ct2$

$\text{dec}(ct, p)$ is just the message, so $\text{dec}(ct, p)$
encrypted under p_2 is the message
encrypted under p_2 . Note that this is a "fresh"
ciphertext, in that the error that was in the
incoming ciphertext does not carry over.

This is difficult to understand, so we can restate it as follows. The decryption procedure $\text{dec}(ct, p)$ is itself a computation, and so it can itself be implemented as a circuit that takes as input the bits of ct and the bits of p , and outputs the decrypted bit $m \in \{0, 1\}$. If someone has a ciphertext ct encrypted under p , a public key for p_2 , and the bits of p encrypted under p_2 , then they can compute $\text{dec}(ct, p) = m$ "homomorphically", and get out m encrypted under p_2 . Notice that the decryption procedure itself washes away the old error; it just outputs 0 or 1.

The decryption procedure is itself a circuit, which contains additions or multiplications, so it will introduce new error, but this new error *does not depend* on the amount of error in the original encryption.

(Note that we can avoid having a distinct new key p_2 (and if you want to bootstrap multiple times, also a p_3 , p_4 ...) by just setting $p_2 = p$. However, this introduces a new assumption, usually called "circular security"; it becomes more difficult to formally prove security if you do this, though many cryptographers think it's fine and circular security poses no significant risk in practice)

But.... there is a catch. In the scheme as described above (using circular security or not), the error blows up so quickly that even the decryption circuit of the scheme itself is too much for it. That is, the new m encrypted under p_2 would *already* have so much error that it is unreadable. This is because each AND gate doubles the bit-length of the error, so a scheme using a d -bit modulus p can only handle less than $\log(d)$ multiplications (in series), but decryption requires computing $\text{mod } p$ in a circuit made up of these binary logic gates, which requires... more than $\log(d)$ multiplications.

Craig Gentry came up with clever techniques to get around this problem, but they are arguably too complicated to explain; instead, we will skip straight to newer work from 2011 and 2013, that solves this problem in a different way.

Learning with errors

To move further, we will introduce a *different type of somewhat-homomorphic encryption* introduced by Brakerski and Vaikuntanathan in 2011, and show how to bootstrap it. Here, we will move away from keys and ciphertexts being *integers*, and instead have keys and ciphertexts be *vectors*. Given a key $k = k_1, k_2, \dots, k_n$, to encrypt a message m , construct a vector $c = c_1, c_2, \dots, c_n$ such that the inner product (or "dot product")

$\langle c, k \rangle = c_1 k_1 + c_2 k_2 + \dots + c_n k_n$, modulo some fixed number p , equals $m + 2e$ where m is the message (which must be 0 or 1), and e is a small (much smaller than p) "error" term. A "public key" that allows encryption but not decryption can be constructed, as before, by making a set of encryptions of 0; an encryptor can randomly combine a subset of these equations and add 1 if the message they are encrypting is 1. To decrypt a ciphertext c knowing the key k , you would compute $\langle c, k \rangle$ modulo p , and see if the result is odd or even (this is the same "mod p mod 2" trick we used earlier). Note that here the $\text{mod } p$ is typically a "symmetric" mod, that is, it returns a number between $-\frac{p}{2}$ and $\frac{p}{2}$ (eg. $137 \bmod 10 = -3$, 212

$\text{mod } 10 = 2$); this allows our error to be positive or negative. Additionally, p does not necessarily have to be prime, though it does need to be odd.

Key	3	14	15	92	65
Ciphertext	2	71	82	81	8

The key and the ciphertext are both vectors, in this example of five elements each.

In this example, we set the modulus $p = 103$. The dot product is $3 * 2 + 14 * 71 + 15 * 82 + 92 * 81 + 65 * 8 = 10202$, and $10202 = 99 * 103 + 5$. 5 itself is of course $2 * 2 + 1$, so the message is 1. Note that in practice, the first element of the key is often set to 1; this makes it easier to generate ciphertexts for a particular value (see if you can figure out why).

The security of the scheme is based on an assumption known as "[learning with errors](#)" (LWE) - or, in more jargony but also more understandable terms, the hardness of *solving systems of equations with errors*.

LWE. The LWE problem asks to recover a secret $\mathbf{s} \in \mathbb{Z}_q^n$ given a sequence of 'approximate' random linear equations on \mathbf{s} . For instance, the input might be

$$\begin{aligned} 14s_1 + 15s_2 + 5s_3 + 2s_4 &\approx 8 \pmod{17} \\ 13s_1 + 14s_2 + 14s_3 + 6s_4 &\approx 16 \pmod{17} \\ 6s_1 + 10s_2 + 13s_3 + 1s_4 &\approx 3 \pmod{17} \\ 10s_1 + 4s_2 + 12s_3 + 16s_4 &\approx 12 \pmod{17} \\ 9s_1 + 5s_2 + 9s_3 + 6s_4 &\approx 9 \pmod{17} \\ 3s_1 + 6s_2 + 4s_3 + 5s_4 &\approx 16 \pmod{17} \end{aligned}$$

where each equation is correct up to some small additive error (say, ± 1), and our goal is to recover \mathbf{s} . (The answer in this case is $\mathbf{s} = (0, 13, 9, 11)$.)

If not for the error, finding \mathbf{s} would be very easy: after about n equations, we can recover \mathbf{s} in polynomial time using Gaussian elimination. Introducing the error seems to make the problem significantly more difficult. For instance, the Gaussian elimination algorithm takes linear combinations of n equations, thereby amplifying the error to unmanageable levels, leaving essentially no information in the resulting equations.

A ciphertext can itself be viewed as an equation: $k_1c_1 + \dots + k_nc_n \approx 0$, where the key $k_1 \dots k_n$ is the unknowns, the ciphertext $c_1 \dots c_n$ is the coefficients, and the equality is only approximate because of both the message (0 or 1) and the error ($2e$ for some relatively small e). The LWE assumption ensures that even if you have access to many of these ciphertexts, you cannot recover k .

Note that in some descriptions of LWE, $\langle c, k \rangle$ can equal any value, but this value must be provided as part of the ciphertext. This is mathematically equivalent to the $\langle c, k \rangle = m+2e$ formulation, because you can just add this answer to the end of the ciphertext and add -1 to the end of the key, and get two vectors that when multiplied together just give $m+2e$. We'll use the formulation that requires $\langle c, k \rangle$ to be near-zero (ie. just $m+2e$) because it is simpler to work with.

Multiplying ciphertexts

It is easy to verify that the encryption is additive: if $\langle ct_1, k \rangle = 2e_1 + m_1$ and $\langle ct_2, k \rangle = 2e_2 + m_2$, then $\langle ct_1 + ct_2, k \rangle = 2(e_1 + e_2) + m_1 + m_2$ (the addition here is modulo p). What is harder is multiplication: unlike with numbers, there is no natural way to multiply two length-n vectors into another length-n vector. The best that we can do is the outer product: a vector containing the products of each possible pair where the first element comes from the first vector and the second element comes from the second vector. That is, $a \otimes b = a_1b_1 + a_2b_1 + \dots + a_nb_1 + a_1b_2 + \dots + a_nb_2 + \dots + a_nb_n$. We can "multiply ciphertexts" using the convenient mathematical identity $\langle a \otimes b, c \otimes d \rangle = \langle a, c \rangle * \langle b, d \rangle$.

Given two ciphertexts c_1 and c_2 , we compute the outer product $c_1 \otimes c_2$. If both c_1 and c_2 were encrypted with k , then $\langle c_1, k \rangle = 2e_1 + m_1$ and $\langle c_2, k \rangle = 2e_2 + m_2$. The outer product $c_1 \otimes c_2$ can be viewed as an encryption of $m_1 * m_2$ under $k \otimes k$; we can see this by looking what happens when we try to decrypt with $k \otimes k$:

$$\begin{aligned} & \langle c_1 \otimes c_2, k \otimes k \rangle \\ &= \langle c_1, k \rangle * \langle c_2, k \rangle \\ &= (2e_1 + m_1) * (2e_2 + m_2) \\ &= 2(e_1m_2 + e_2m_1 + 2e_1e_2) + m_1m_2 \end{aligned}$$

So this outer-product approach works. But there is, as you may have already noticed, a catch: the size of the ciphertext, and the key, grows quadratically.

Relinearization

We solve this with a **relinearization** procedure. The holder of the private key \mathbf{k} provides, as part of the public key, a "relinearization key", which you can think of as "noisy" encryptions of $\mathbf{k} \otimes \mathbf{k}$ under \mathbf{k} . The idea is that we provide these encrypted pieces of $\mathbf{k} \otimes \mathbf{k}$ to anyone performing the computations, allowing them to compute the equation $\langle \mathbf{c}_1 \otimes \mathbf{c}_2, \mathbf{k} \otimes \mathbf{k} \rangle$ to "decrypt" the ciphertext, but only in such a way that the output comes back encrypted under \mathbf{k} .

It's important to understand what we mean here by "noisy encryptions". Normally, this encryption scheme only allows encrypting $\mathbf{m} \in \{0, 1\}$, and an "encryption of \mathbf{m} " is a vector \mathbf{c} such that $\langle \mathbf{c}, \mathbf{k} \rangle = \mathbf{m} + 2e$ for some small error e . Here, we're "encrypting" arbitrary $\mathbf{m} \in \{0, 1, 2, \dots, p-1\}$. Note that the error means that you can't fully recover \mathbf{m} from \mathbf{c} ; your answer will be off by some multiple of 2. However, it turns out that, for this specific use case, this is fine.

The relinearization key consists of a set of vectors which, when inner-producted (modulo p) with the key \mathbf{k} , give values of the form $\mathbf{k}_i * \mathbf{k}_j * 2^d + 2e \pmod{p}$, one such vector for every possible triple (i, j, d) , where i and j are indices in the key and d is an exponent where $2^d < p$ (note: if the key has length n , there would be $n^2 * \log(p)$ values in the relinearization key; make sure you understand why before continuing).

enc($\mathbf{k}_1 * \mathbf{k}_1$)	enc($\mathbf{k}_1 * \mathbf{k}_1 * 2$)	enc($\mathbf{k}_1 * \mathbf{k}_1 * 4$)	enc($\mathbf{k}_1 * \mathbf{k}_1 * 8$)
enc($\mathbf{k}_1 * \mathbf{k}_2$)	enc($\mathbf{k}_1 * \mathbf{k}_2 * 2$)	enc($\mathbf{k}_1 * \mathbf{k}_2 * 4$)	enc($\mathbf{k}_1 * \mathbf{k}_2 * 8$)
enc($\mathbf{k}_2 * \mathbf{k}_1$)	enc($\mathbf{k}_2 * \mathbf{k}_1 * 2$)	enc($\mathbf{k}_2 * \mathbf{k}_1 * 4$)	enc($\mathbf{k}_2 * \mathbf{k}_1 * 8$)
enc($\mathbf{k}_2 * \mathbf{k}_2$)	enc($\mathbf{k}_2 * \mathbf{k}_2 * 2$)	enc($\mathbf{k}_2 * \mathbf{k}_2 * 4$)	enc($\mathbf{k}_2 * \mathbf{k}_2 * 8$)

Example assuming $p = 15$ and \mathbf{k} has length 2. Formally, $\text{enc}(x)$ here means "outputs $x+2e$ if inner-producted with \mathbf{k} ".

Now, let us take a step back and look again at our goal. We have a ciphertext which, if decrypted with $\mathbf{k} \otimes \mathbf{k}$, gives $\mathbf{m}_1 * \mathbf{m}_2$. We want a ciphertext which, if decrypted with \mathbf{k} , gives $\mathbf{m}_1 * \mathbf{m}_2$. We can do this with the relinearization key. Notice that the decryption equation $\langle \mathbf{c}_1 \otimes \mathbf{c}_2, \mathbf{k} \otimes \mathbf{k} \rangle$ is just a big sum of terms of the form $(\mathbf{c}_{1,i} * \mathbf{c}_{2,j}) * \mathbf{k}_p * \mathbf{k}_q$.

And what do we have in our relinearization key? A bunch of elements of the form $2^d * k_p * k_q$, noisy-encrypted under \mathbf{k} , for every possible combination of p and q ! Having all the powers of two in our relinearization key allows us to generate any $(ct_{1,i} * ct_{2,j}) * k_p * k_q$ by just adding up $\leq \log(p)$ powers of two (eg. $13 = 8 + 4 + 1$) together for each (p, q) pair.

For example, if $ct_1 = [1, 2]$ and $ct_2 = [3, 4]$, then $ct_1 \otimes ct_2 = [3, 4, 6, 8]$, and $enc(<ct_1 \otimes ct_2, k \otimes k>) = enc(3k_1k_1 + 4k_1k_2 + 6k_2k_1 + 8k_2k_2)$ could be computed via:

$$enc(k_1 * k_1) + enc(k_1 * k_1 * 2) + enc(k_1 * k_2 * 4) + \\ enc(k_2 * k_1 * 2) + enc(k_2 * k_1 * 4) + enc(k_2 * k_2 * 8)$$

Note that each noisy-encryption in the relinearization key has some even error $2e$, and the equation $<ct_1 \otimes ct_2, k \otimes k>$ itself has some error: if $<ct_1, k> = 2e_1 + m_1$ and $<ct_2, k> = 2e_2 + m_2$, then $<ct_1 \otimes ct_2, k \otimes k> = <ct_1, k> * <ct_2, k> = 2(2e_1e_2 + e_1m_2 + e_2m_1) + m_1m_2$. But this total error is still (relatively) small ($2e_1e_2 + e_1m_2 + e_2m_1$ plus $n^2 * \log(p)$ fixed-size errors from the relinearization key), and the error is even, and so the result of this calculation still gives a value which, when inner-producted with \mathbf{k} , gives $m_1 * m_2 + 2e'$ for some "combined error" e' .

The broader technique we used here is a common trick in homomorphic encryption: provide pieces of the key encrypted under the key itself (or a different key if you are pedantic about avoiding circular security assumptions), such that someone computing on the data can compute the decryption equation, but only in such a way that the output itself is still encrypted. It was used in bootstrapping above, and it's used here; it's best to make sure you mentally understand what's going on in both cases.

This new ciphertext has considerably more error in it: the $n^2 * \log(p)$ different errors from the portions of the relinearization key that we used, plus the $2(2e_1e_2 + e_1m_2 + e_2m_1)$ from the original outer-product ciphertext. Hence, the new ciphertext still does have quadratically larger error than the original ciphertexts, and so we still haven't solved the problem that the error blows up too quickly. To solve this, we move on to another trick...

Modulus switching

Here, we need to understand an important algebraic fact. A ciphertext is a vector \mathbf{ct} , such that $<\mathbf{ct}, \mathbf{k}> = \mathbf{m} + 2\mathbf{e}$, where $\mathbf{m} \in \{0, 1\}$. But we can also look at the ciphertext from a different "perspective": consider $\frac{\mathbf{ct}}{2}$ (modulo p). $<\frac{\mathbf{ct}}{2}, \mathbf{k}> = \frac{\mathbf{m}}{2} + \mathbf{e}$, where $\frac{\mathbf{m}}{2} \in \{0, \frac{p+1}{2}\}$. Note that

because (modulo p) $(\frac{p+1}{2}) * 2 = p + 1 = 1$, division by 2 (modulo p) maps 1 to $\frac{p+1}{2}$; this is a very convenient fact for us.

x	Modular division by 2, mod 9	Regular rounded-down division by 2
0	0	0
1	5	0
2	1	1
3	6	1
4	2	2
5	7	2
6	3	3
7	8	3
8	4	4

The scheme in this section uses both modular division (ie. multiplying by the modular multiplicative inverse) and regular "rounded down" integer division; make sure you understand how both work and how they are different from each other.

That is, the operation of dividing by 2 (modulo p) converts small even numbers into small numbers, and it converts 1 into $\frac{p}{2}$ (rounded up). So if we look at $\frac{ct}{2}$ (modulo p) instead of ct , decryption involves computing $\langle \frac{ct}{2}, k \rangle$ and seeing if it's closer to 0 or $\frac{p}{2}$. This "perspective" is much more robust to certain kinds of errors, where you know the error is small but can't guarantee that it's a multiple of 2.

Now, here is something we can do to a ciphertext.

1. Start: $\langle ct, k \rangle = \{0 \text{ or } 1\} + 2e \pmod{p}$
2. Divide ct by 2 (modulo p): $\langle ct', k \rangle = \{0 \text{ or } \frac{p}{2}\} + e \pmod{p}$
3. Multiply ct' by $\frac{q}{p}$ using "regular rounded-down integer division":
 $\langle ct'', k \rangle = \{0 \text{ or } \frac{q}{2}\} + e' + e_2 \pmod{q}$
4. Multiply ct'' by 2 (modulo q): $\langle ct''', k \rangle = \{0 \text{ or } 1\} + 2e' + 2e_2 \pmod{q}$

Step 3 is the crucial one: it converts a ciphertext under modulus p into a ciphertext under modulus q . The process just involves "scaling down" each element of \mathbf{ct}' by multiplying by $\frac{q}{p}$ and rounding down, eg. $\text{floor}(56 * \frac{15}{103}) = \text{floor}(8.15533\ldots) = 8$.

The idea is this: if $\langle \mathbf{ct}', k \rangle = m * \frac{p}{2} + e \pmod{p}$, then we can interpret this as

$$\langle \mathbf{ct}', k \rangle = p(z + \frac{m}{2}) + e \text{ for some integer } z. \text{ Therefore, } \langle \mathbf{ct}' * \frac{q}{p}, k \rangle = q(z + \frac{m}{2}) + e * \frac{p}{q}.$$

Rounding adds error, but only a little bit (specifically, up to the size of the values in k , and we can make the values in k small without sacrificing security). Therefore, we can say

$$\langle \mathbf{ct}' * \frac{q}{p}, k \rangle = m * \frac{q}{2} + e' + e_2 \pmod{q}, \text{ where } e' = e * \frac{q}{p}, \text{ and } e_2 \text{ is a small error from rounding.}$$

What have we accomplished? We turned a ciphertext with modulus p and error $2e$ into a ciphertext with modulus q and error $2(\text{floor}(e * \frac{p}{q}) + e_2)$, where the new error is *smaller* than the original error.

Let's go through the above with an example. Suppose:

- \mathbf{ct} is just one value, [5612]
- $k = [9]$
- $p = 9999$ and $q = 113$

$\langle \mathbf{ct}, k \rangle = 5612 * 9 = 50508 = 9999 * 5 + 2 * 256 + 1$, so \mathbf{ct} represents the bit 1, but the error is fairly large ($e = 256$).

Step 2: $\mathbf{ct}' = \frac{\mathbf{ct}}{2} = 2806$ (remember this is modular division; if \mathbf{ct} were instead 5613, then we would have $\frac{\mathbf{ct}}{2} = 7806$). Checking: $\langle \mathbf{ct}', k \rangle = 2806 * 9 = 25254 = 9999 * 2.5 + 256.5$

Step 3: $\mathbf{ct}'' = \text{floor}(2806 * \frac{113}{9999}) = \text{floor}(31.7109\ldots) = 31$. Checking:

$$\langle \mathbf{ct}'', k \rangle = 279 = 113 * 2.5 - 3.5$$

Step 4: $\mathbf{ct}''' = 31 * 2 = 62$. Checking: $\langle \mathbf{ct}''', k \rangle = 558 = 113 * 5 - 2 * 4 + 1$

And so the bit 1 is preserved through the transformation. The crazy thing about this procedure is: *none of it requires knowing k* . Now, an astute reader might notice: you reduced the *absolute* size of the error (from 256 to 2), but the *relative* size of the error remained unchanged, and even slightly increased: $\frac{256}{9999} \approx 2.5\%$ but $\frac{4}{113} \approx 3.5\%$. Given that it's the relative error that causes ciphertexts to break, what have we gained here?

The answer comes from what happens to error when you multiply ciphertexts. Suppose that we start with a ciphertext x with error 100, and modulus $p = 10^{16} - 1$. We want to repeatedly square x , to compute $((x^2)^2)^2 = x^{16}$. First, the "normal way":

Power of x	Error	Modulus
x	100	$10^{16}-1$
x^2	10^4	$10^{16}-1$
x^4	10^8	$10^{16}-1$
x^8	10^{16}	$10^{16}-1$
x^{16}	10^{32}	$10^{16}-1$

The error blows up too quickly for the computation to be possible. Now, let's do a modulus reduction after every multiplication. We assume the modulus reduction is imperfect and increases error by a factor of 10, so a 1000x modulo reduction only reduces error from 10000 to 100 (and not to 10):

Power of x	Error	Modulus
x	100	$10^{16}-1$
x^2	10^4	$10^{16}-1$
x^4	100	$10^{13}-1$
x^8	10^4	$10^{13}-1$
x^{16}	100	$10^{10}-1$
x^8	10^4	$10^{10}-1$
x^4	100	10^7-1
x^2	10^4	10^7-1
x	100	$10^{16}-1$

The key mathematical idea here is that the *factor* by which error increases in a multiplication depends on the absolute size of the error, and not its relative size, and so if we keep doing modulus reductions to keep the error small, each multiplication only increases the error by a

constant factor. And so, with a d bit modulus (and hence $\approx 2^d$ room for "error"), we can do $O(d)$ multiplications! This is enough to bootstrap.

Another technique: matrices

Another technique (see Gentry, Sahai, Waters (2013)) for fully homomorphic encryption involves matrices: instead of representing a ciphertext as ct where $\langle ct, k \rangle = 2e + m$, a ciphertext is a matrix, where $k * CT = k * m + e$ (k , the key, is still a vector). The idea here is that k is a "secret near-eigenvector" - a secret vector which, if you multiply the matrix by it, returns something very close to either zero or the key itself.

The fact that addition works is easy: if $k * CT_1 = m_1 * k + e_1$ and $k * CT_2 = m_2 * k + e_2$, then $k * (CT_1 + CT_2) = (m_1 + m_2) * k + (e_1 + e_2)$. The fact that multiplication works is also easy:

$$\begin{aligned} k * CT_1 * CT_2 &= (m_1 * k + e_1) * CT_2 = m_1 * k * CT_2 + e_1 * CT_2 \\ &= m_1 * m_2 * k + m_1 * e_2 + e_1 * CT_2 \end{aligned}$$

The first term is the "intended term"; the latter two terms are the "error". That said, notice that here error does blow up quadratically (see the $e_1 * CT_2$ term; the size of the error increases by the size of each ciphertext element, and the ciphertext elements also square in size), and you do need some clever tricks for avoiding this. Basically, this involves turning ciphertexts into matrices containing their constituent bits before multiplying, to avoid multiplying by anything higher than 1; if you want to see how this works in detail I recommend looking at my code:
https://github.com/vbuterin/research/blob/master/matrix_fhe/matrix_fhe.py#L121

In addition, the code there, and also

https://github.com/vbuterin/research/blob/master/tensor_fhe/homomorphic_encryption.py#L186, provides simple examples of useful circuits that you can build out of these binary logical operations; the main example is for adding numbers that are represented as multiple bits, but one can also make circuits for comparison ($<$, $>$, $=$), multiplication, division, and many other operations.

Since 2012-13, when these algorithms were created, there have been many optimizations, but they all work on top of these basic frameworks. Often, polynomials are used instead of integers; this is called [ring LWE](#). The major challenge is still efficiency: an operation involving a single bit involves multiplying entire matrices or performing an entire relinearization

computation, a very high overhead. There are tricks that allow you to perform many bit operations in a single ciphertext operation, and this is actively being worked on and improved.

We are quickly getting to the point where many of the applications of homomorphic encryption in privacy-preserving computation are starting to become practical. Additionally, research in the more advanced applications of the lattice-based cryptography used in homomorphic encryption is rapidly progressing. So this is a space where some things can already be done today, but we can hopefully look forward to much more becoming possible over the next decade.